



Efficient and Customizable Dynamic Hand Gesture Recognition for Edge Deployment



Victor Teslaru; Gabriel Pojoga
"Gheorghe Asachi" Technical University of Iasi

1. Problem & Scope

Challenges:

- Real-time performance on resource-constrained devices
- Domain-specific vocabularies need custom datasets, no universal solution
- Ambiguous user intent from natural movement hard to capture in the collection phase

Our solution:

- Hierarchical pipeline (palm → landmarks → gesture)
- Active learning, guides data collection to fix weaknesses

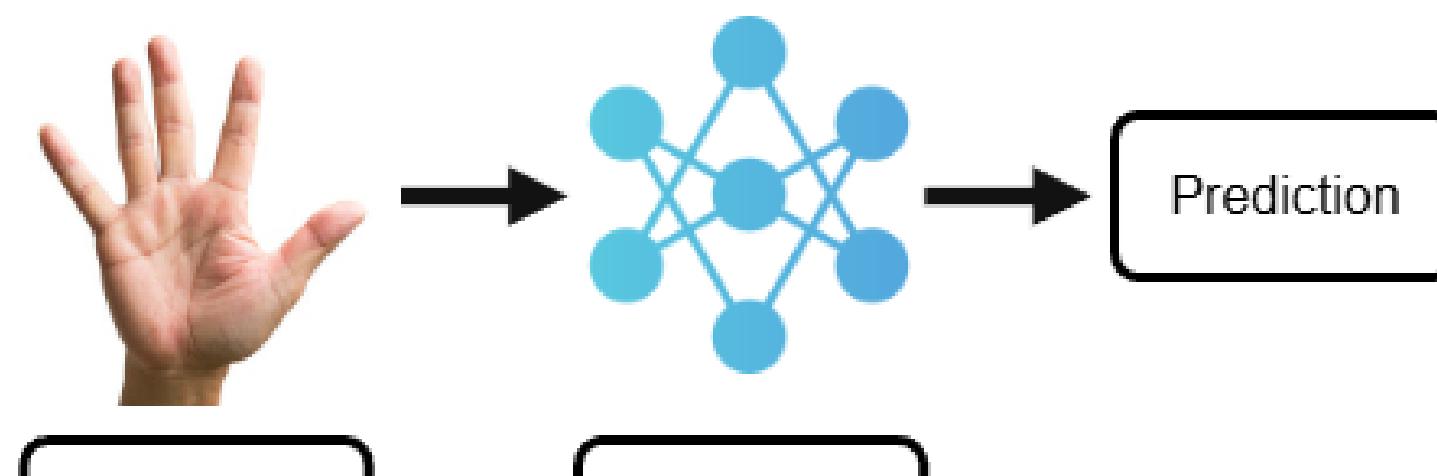


Figure 1. End-to-end approach.

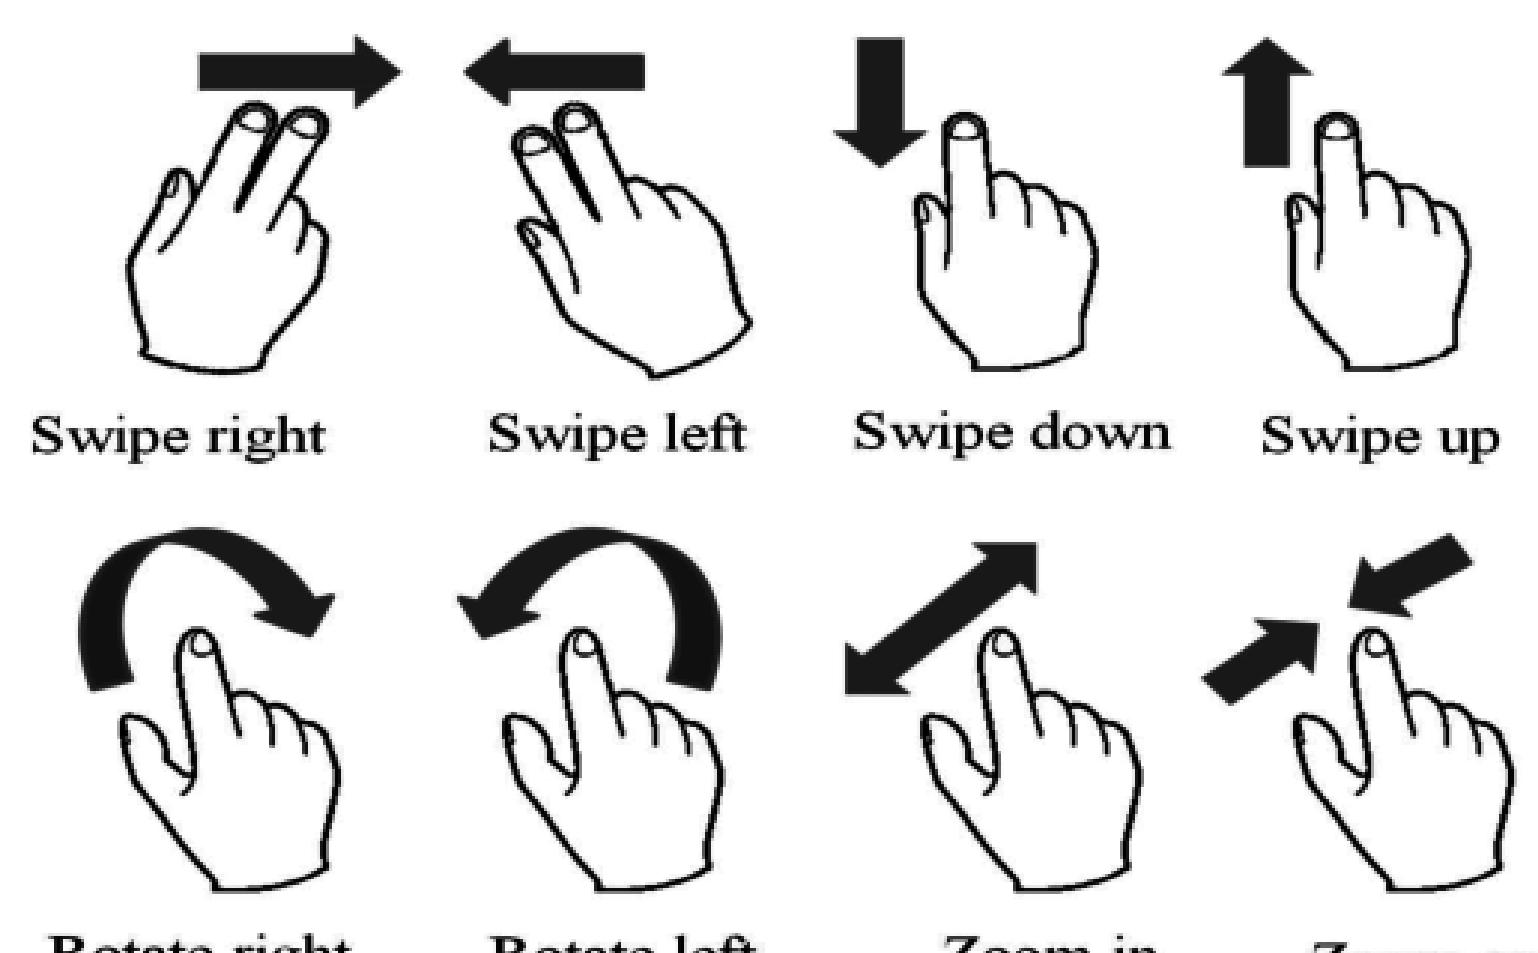


Figure 2. Dynamic hand gestures captured (taken from [2]).

2. Data Pipeline & Dataset

Collection:

- Dynamic Time Warping [1] similarity to prevent redundancy
- Interactive pose rotation augmentation with user validation
- Linear interpolation for ≤ 2 missing frames, discard otherwise

Dataset:

- 8 gestures: Swipe (4), Zoom (2), Rotate (2)
- 10-frame sequences, 21 3D landmarks
- Wrist and hand span normalized (position/scale-invariant)

3. Model Design & Edge Optimization

Hierarchical architecture:

- SSD-based palm detector [3], which reduces search space
- 21 key-points extraction for structured representation
- NPU offloading for vision tasks
- Quantized 1D CNN for temporal modelling [4]

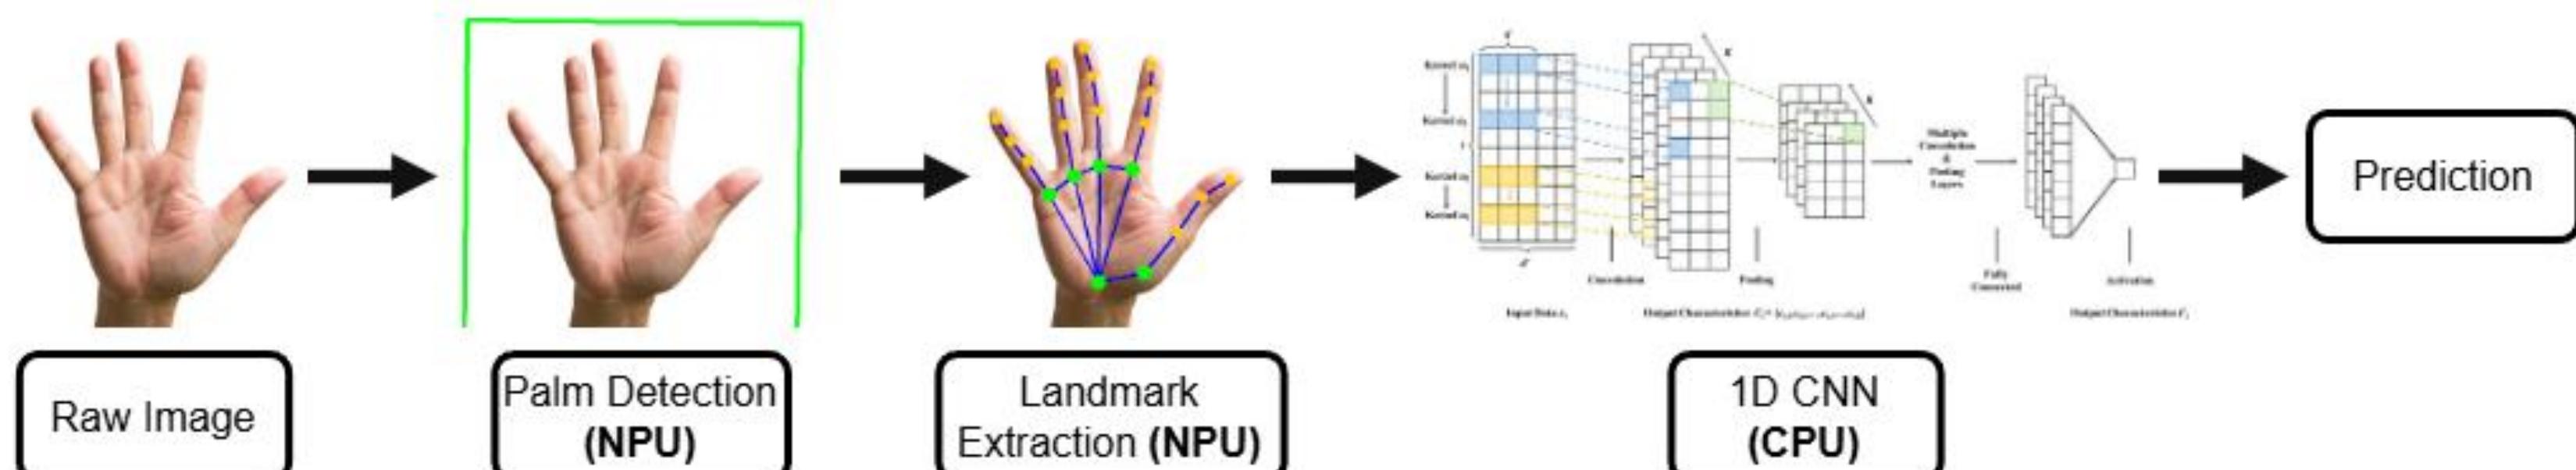


Figure 3. Hierarchical architecture pipeline.

4. Deployment & Inference

Sliding window continuous inference:

- Capture 10-20 frame sequences (user-paced)
- Extract up to 7 overlapping inverse sliding windows, prioritizing recency
- Classify each window independently using uncertainty thresholding
- Filter out ambiguities through majority voting across valid predictions

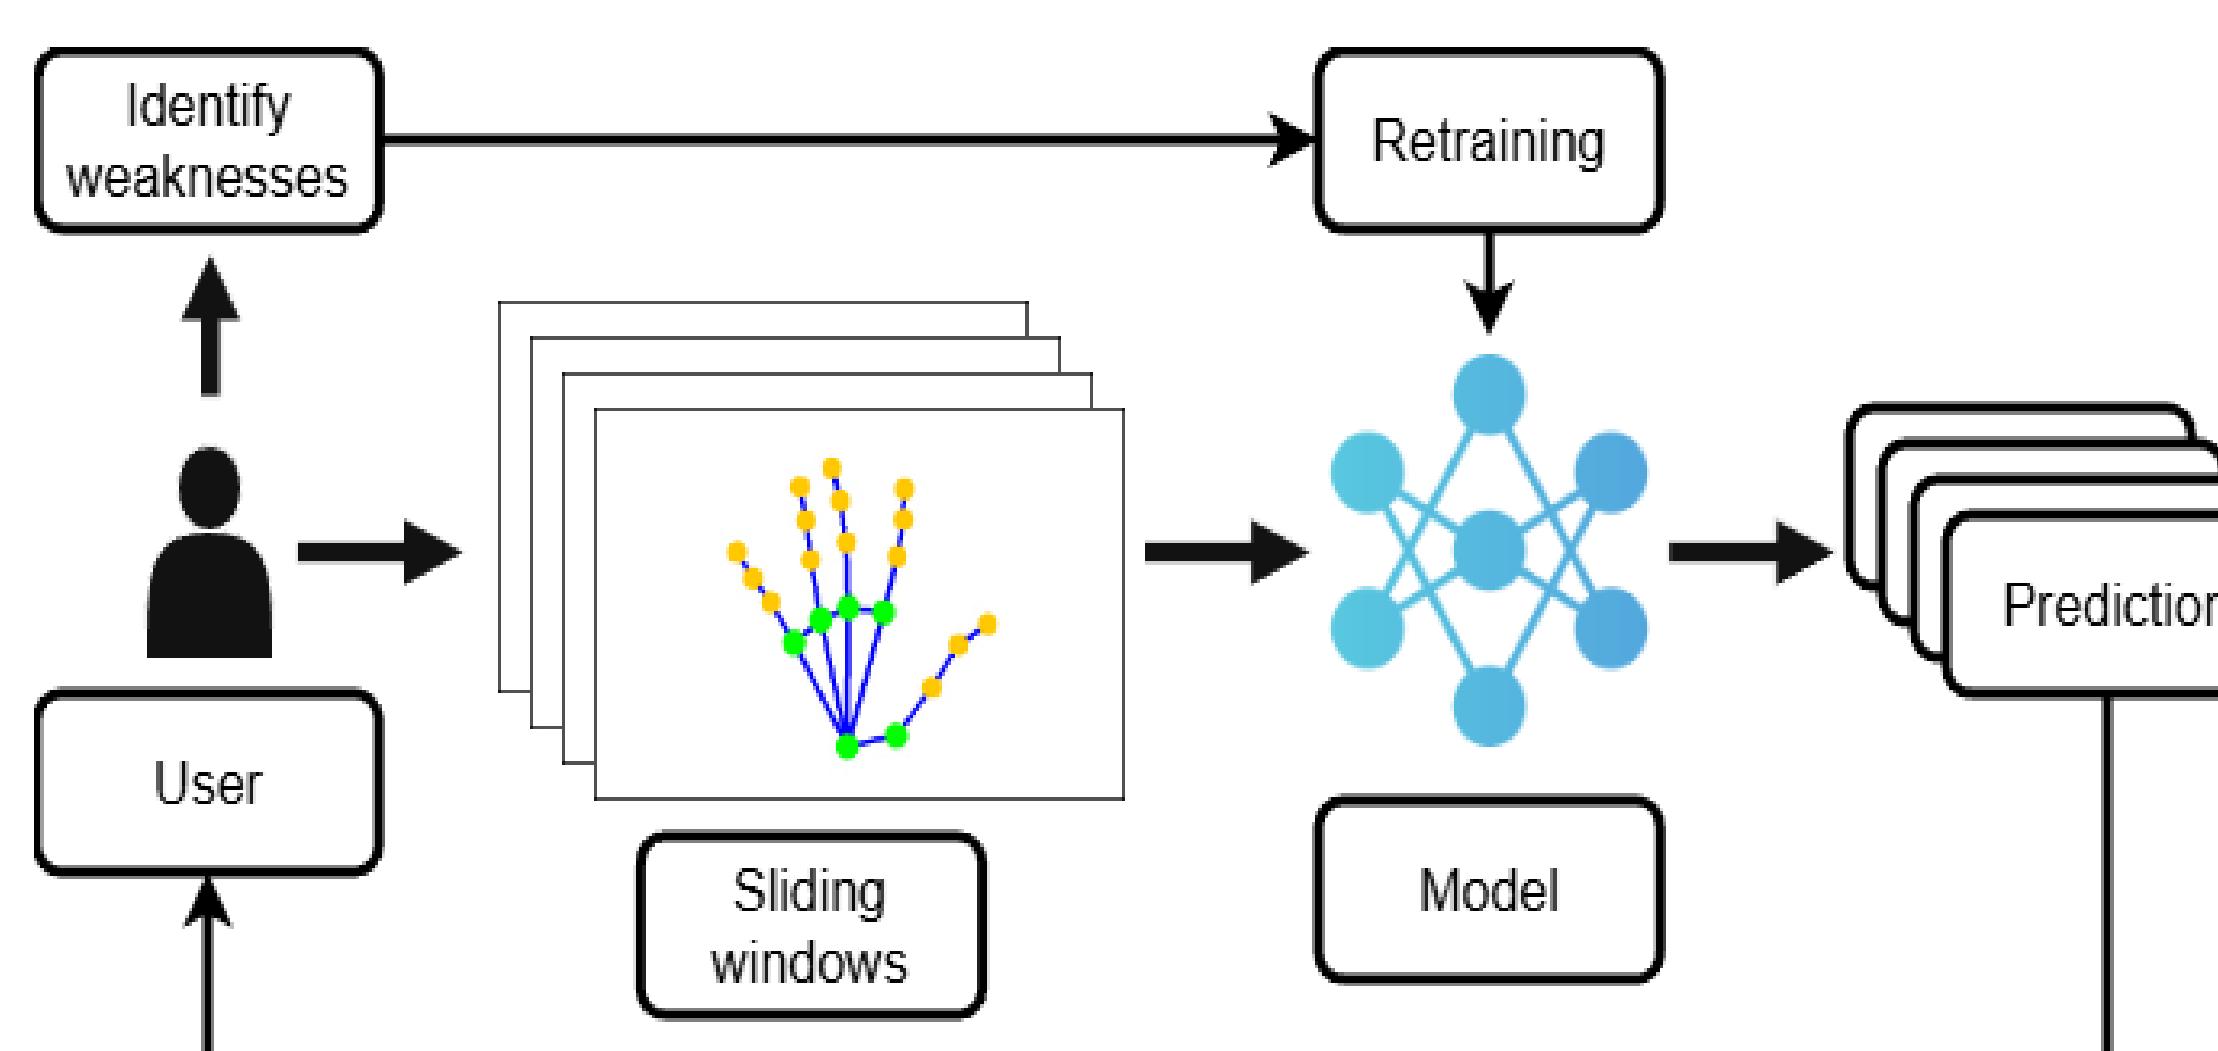


Figure 4. Active learning workflow.

5. Active Learning

Failure driven improvement:

- Actual failure points from source (not hypothetical)
- Multi-window anomalies caught automatically alongside natural movement
- Can adapt to new user variation

Workflow:

- Multiple window predictions
- User reviews (labels / discards)
- Diversity check + augmentation
- Incremental dataset update

6. Evaluation metrics

- Benchmark: 73 gestures with various poses and sliding windows length
- Baseline struggles with unknown class based only on thresholding
- Active learning helps the model better define decision boundaries
- Unknown gesture recall: 20% → 95% (4.75x improvement)

| Baseline Model | | | | | | | | | | Enhanced Model | | | | | | | | | |
|----------------|---|---|---|---|---|----|---|---|---|----------------|---|---|---|---|---|---|---|---|----|
| | | | | | | | | | | | | | | | | | | | |
| True | | | | | | | | | | Predicted | | | | | | | | | |
| Swipe Up - | 8 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | Swipe Up - | 8 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| Swipe Down - | 1 | 3 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | Swipe Down - | 0 | 5 | 0 | 0 | 0 | 0 | 0 | 0 | 1 |
| Swipe Left - | 0 | 0 | 4 | 0 | 0 | 0 | 0 | 0 | 2 | Swipe Left - | 0 | 0 | 7 | 0 | 0 | 0 | 0 | 0 | 0 |
| Swipe Right - | 0 | 0 | 0 | 6 | 0 | 0 | 1 | 0 | 0 | Swipe Right - | 0 | 0 | 0 | 6 | 0 | 0 | 0 | 0 | 1 |
| Zoom In - | 2 | 0 | 0 | 0 | 5 | 0 | 1 | 0 | 0 | Zoom In - | 0 | 0 | 0 | 0 | 6 | 0 | 0 | 0 | 2 |
| Zoom Out - | 1 | 0 | 0 | 0 | 1 | 4 | 0 | 0 | 0 | Zoom Out - | 0 | 0 | 0 | 0 | 0 | 6 | 0 | 0 | 0 |
| Rotate Clo - | 0 | 0 | 0 | 0 | 0 | 0 | 4 | 0 | 1 | Rotate Clo - | 0 | 0 | 0 | 0 | 0 | 0 | 5 | 0 | 0 |
| Rotate Cou - | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 5 | 0 | Rotate Cou - | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 6 | 0 |
| Unknown - | 1 | 0 | 0 | 1 | 1 | 11 | 0 | 2 | 4 | Unknown - | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 19 |

Figure 5. Model evaluation before and after active learning.

Contact

Victor Teslaru
Gabriel Pojoga
"Gheorghe Asachi" Technical University of Iasi
Email: victor.teslaru@student.tuiasi.ro
gabriel.pojoga@student.tuiasi.ro

References

1. H. Sakoe and S. Chiba, "Dynamic programming algorithm optimization for spoken word recognition," in IEEE Transactions on Acoustics, Speech, and Signal Processing, vol. 26, no. 1, pp. 43-49, February 1978, doi: 10.1109/TASSP.1978.1163055.
2. Tumuganti, Nagarkarthy & Ahn, Eun & Bae, Yun & Choi, Jun. (2013). TCAM based pattern matching technique for hand gesture recognition. 368-369. 10.1109/ISOCC.2013.6864052.
3. Zhang, F., Bazarevsky, V., Vakunov, A., Tkachenko, A., Sung, G., Chang, C. L., & Grundmann, M. (2020). Mediapipe hands: On-device real-time hand tracking. arXiv preprint arXiv:2006.10214.
4. TensorFlow Team. "Post-training quantization," TensorFlow Lite Documentation, 2023. https://www.tensorflow.org/lite/performance/post_training_quantization